Output representations

Inputs and outputs of networks

$f_\theta : X \to O$

- Input:
  - Tensor $\mathbf{x}$
- Output:
  - Tensor $\mathbf{o}$

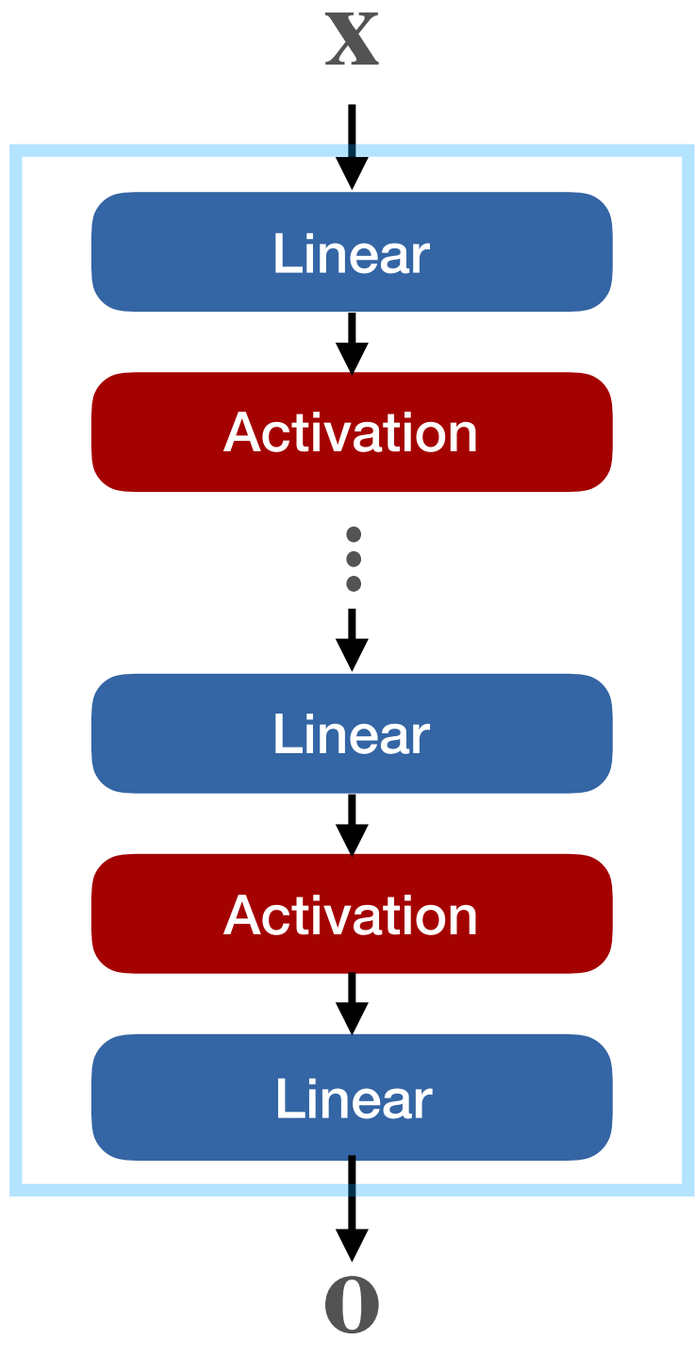

Regression

- vanilla tensor $\hat{\mathbf{y}} = \mathbf{0}$

Positive regression

- Option 1: ReLU
  - $\hat{\mathbf{y}} = \max(\mathbf{o}, 0)$

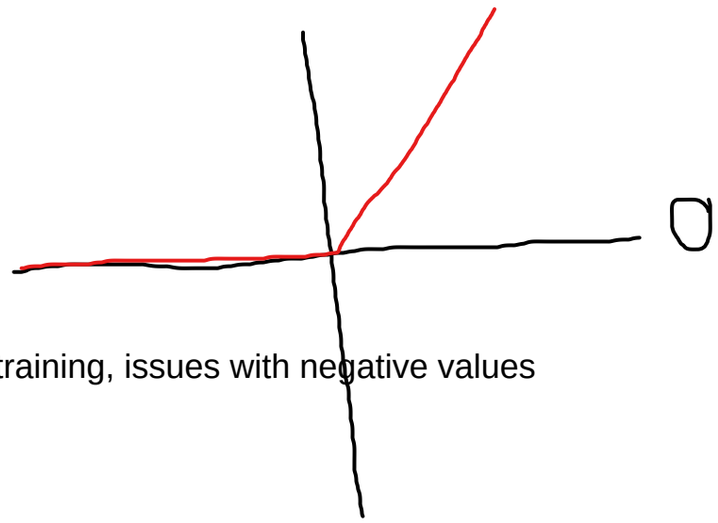

Don't use ReLU for regression training, issues with negative values

- Option 2: Soft ReLU
  - $\hat{\mathbf{y}} = \log(1 + e^{\mathbf{o}})$

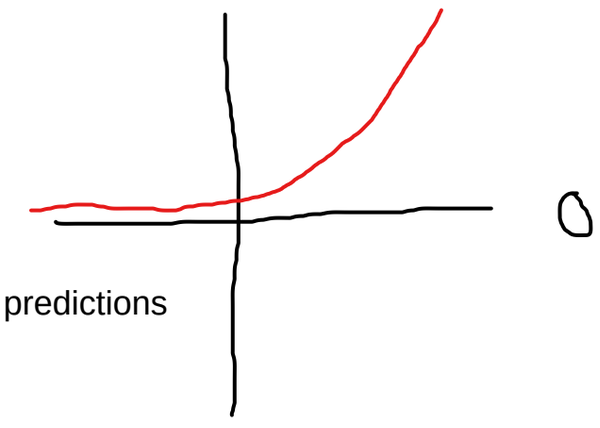

This function can recover from negative predictions

Binary Classification

- Option 1: Thresholding
  - $\hat{\mathbf{y}} = \mathbf{o} > 0$

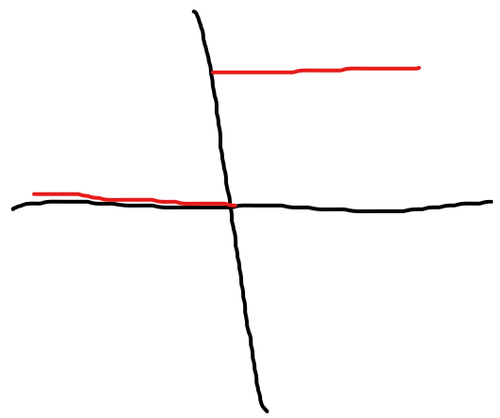

hard to train, the gradient itself of this function is going to be 0 everywhere.

Does not give you a signal on how to fit the network better

- Option 2: Logistic Regression
  - $p(1) = \sigma(\mathbf{o})$

General Classification

- Output more values, one per class
- Option 1: argmax
  - $\hat{y} = \text{argmax}_i \mathbf{o}_i$

can not train the network, can not compute the gradient with this function

- Option 2: softmax
  - $p(y) = \text{softmax}(\mathbf{o})_y$

Output representations in practice

- Do not add into model
- Always output raw values

Never add an output transformation inside your network because output transformations are hard to differentiate through.