Activation functions

Non-linearities

- Allow a deep network to model arbitrary differentiable functions

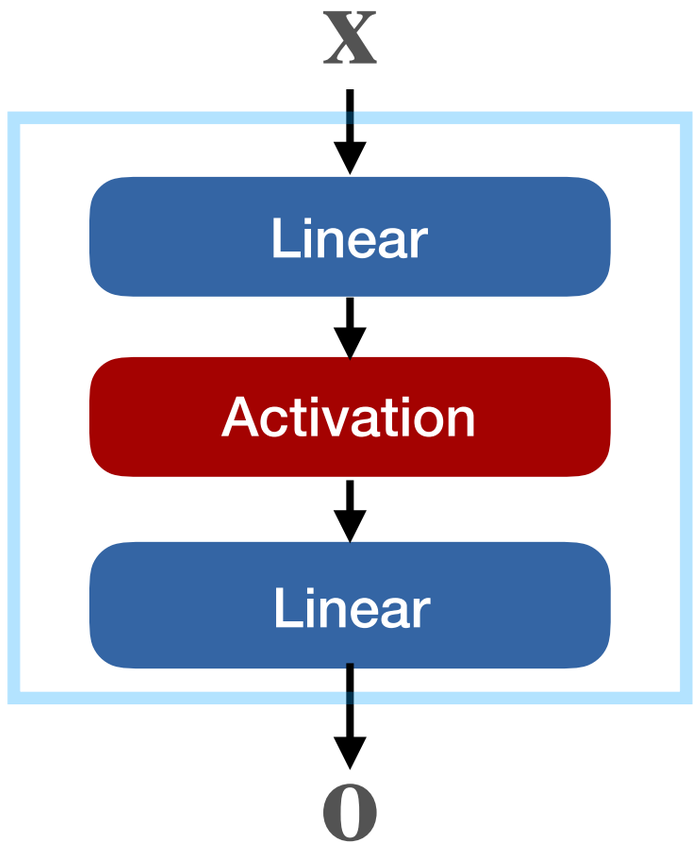

Zoo of activation functions

ReLU

Leaky ReLU

PReLU

ELU

tanh

Sigmoid

Maxout

Sigmoid

- $\frac{1}{1+e^{-x}}$
- Same as tanh
- Do not use

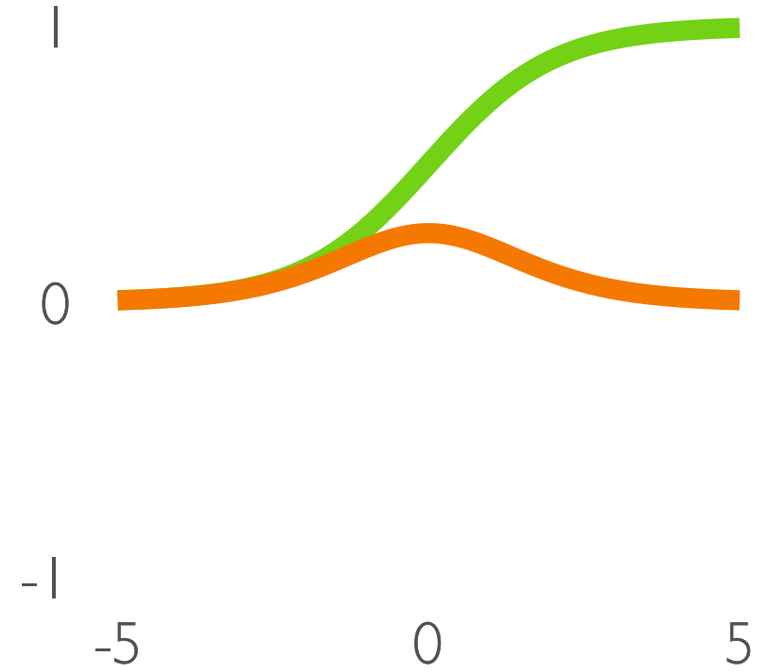

ReLU

- max(x,0)

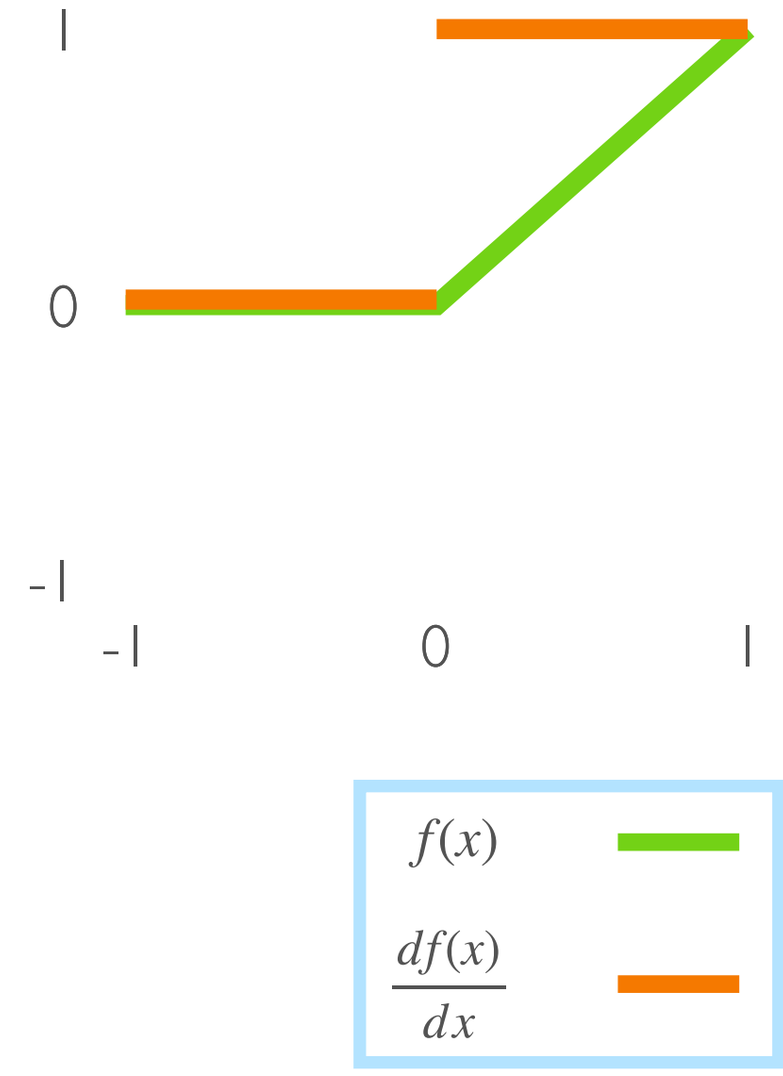

Dead ReLUs

- Prevent dead ReLUs:
  - Initialize Network carefully
  - Decrease the learning rate

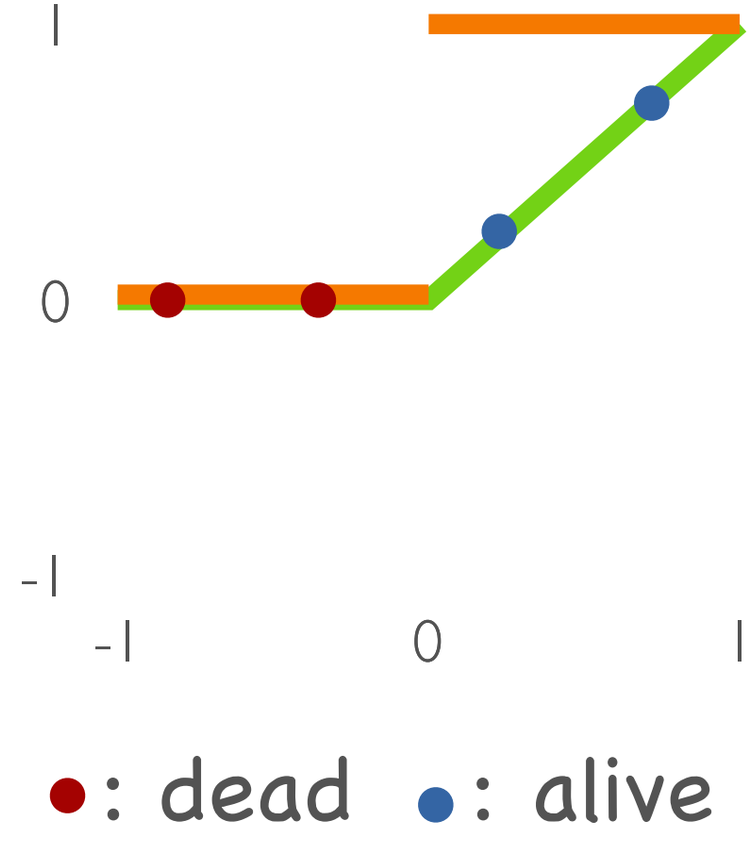

Leaky ReLU

- $\max(x, \alpha x)$
- For $0 < \alpha < 1$
- Called PReLU if α is learned

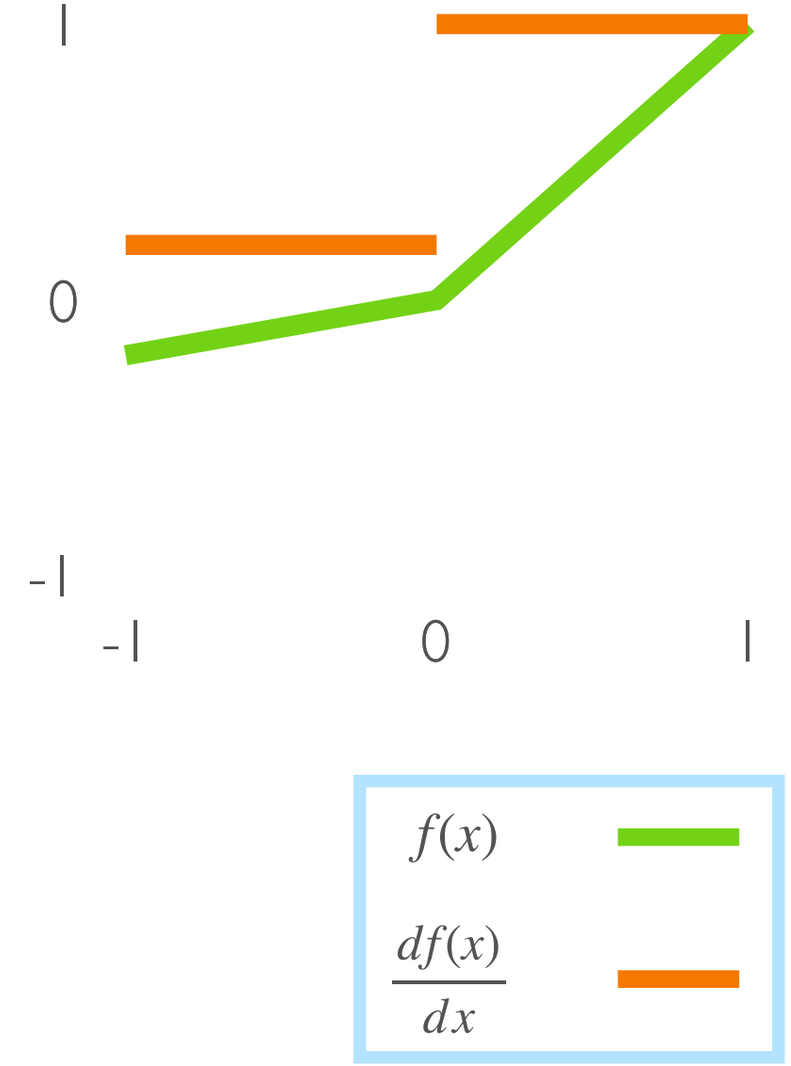

ELU

ELU more expensive to compute

- $$\begin{cases} x & \text{for } x \geq 0 \\ \alpha(e^x - 1) & \text{for } x < 0 \end{cases}$$

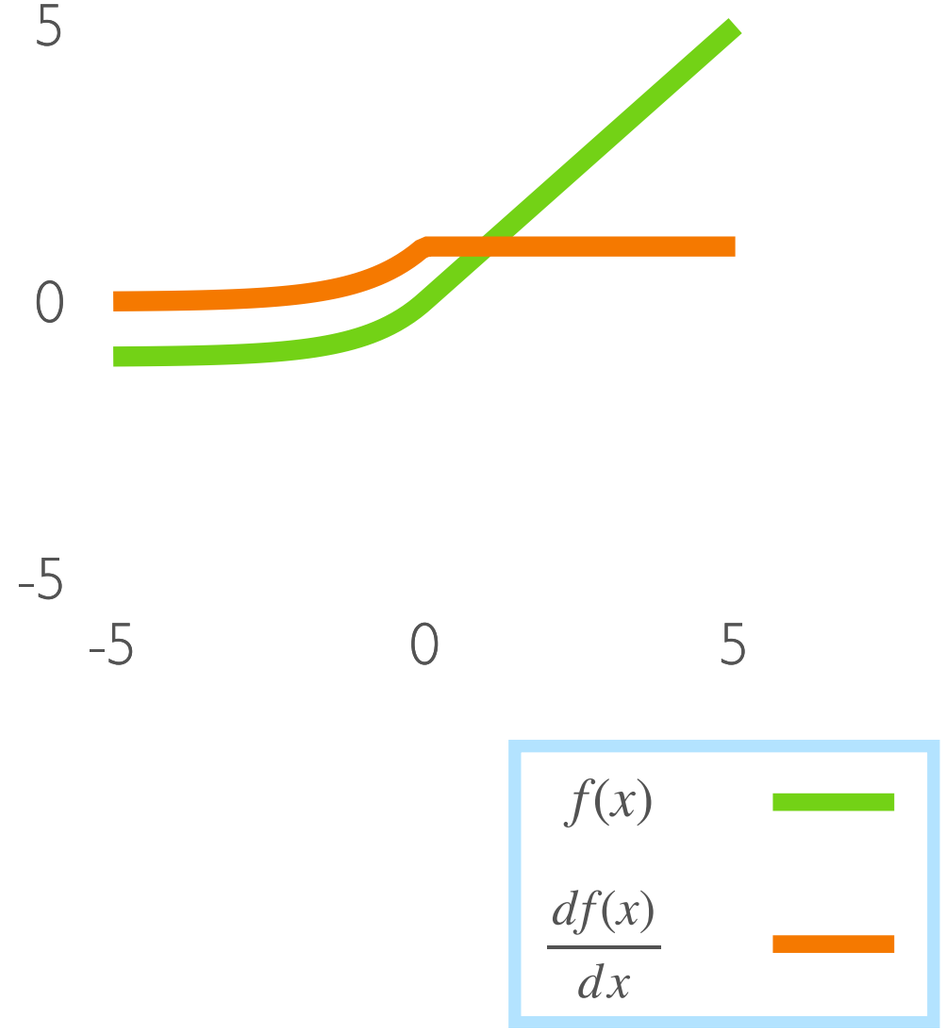

Maxout

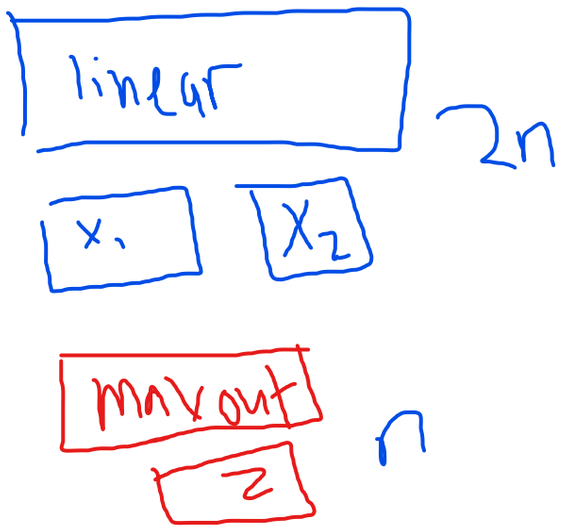

- $\max(x_1, x_2)$

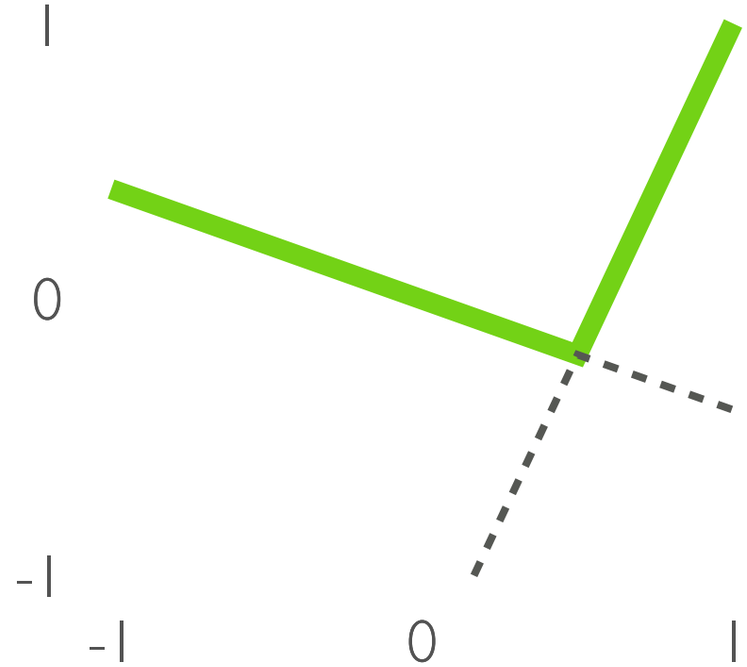

Which activation to choose?

- ReLU
  - Carefully initialize
    - Small enough learning rate
- If ReLU fails, try:
  - Leaky ReLU / PReLU
- Avoid sigmoid and tanh